


Our CNN (GEN4, native HST) vs. TinyLensGPU (Cao et al. 2025) — same 63 SLACS, Bolton 2008 GT

Predicted vs. true θ_E (63 SLACS)



Error distribution (lower-right = more accurate)

